Average Cumulative Regret with Upper Confidence Intervals - Bernoulli Environment

Legend:
- EpsilonGreedy(epsilon=0.1, bernoulli)
- UCB
- KL-UCB
- ThompsonSampling(bernoulli)
- GradientBandit(α=0.1, with_baseline)
- LLM (gpt-4.1-nano)
- LLM(gpt-4.1-nano)

X-axis: Steps
Y-axis: Cumulative Regret